

ICLR 2015

# ADAM: A Method for Stochastic Optimization

Diederik P. Kingma \*

University of Amsterdam, OpenAI

Jimmy Lei Ba \*

University of Toronto

\* Equal contribution

# What is optimization?

The process of finding the best solution from all possible solutions

# What is optimization in machine learning?

The process of adjusting model parameters to minimize the objective function (loss function) and improve model performance.

# Line-up

1

2

3

4

5

6

7

Optimization  
algorithms

Related  
work

Adam

Convergence  
analysis

Codes

Experiments

Limitations

# Line-up

1

2

3

4

5

6

7

Optimization  
algorithms

Related  
work

Adam

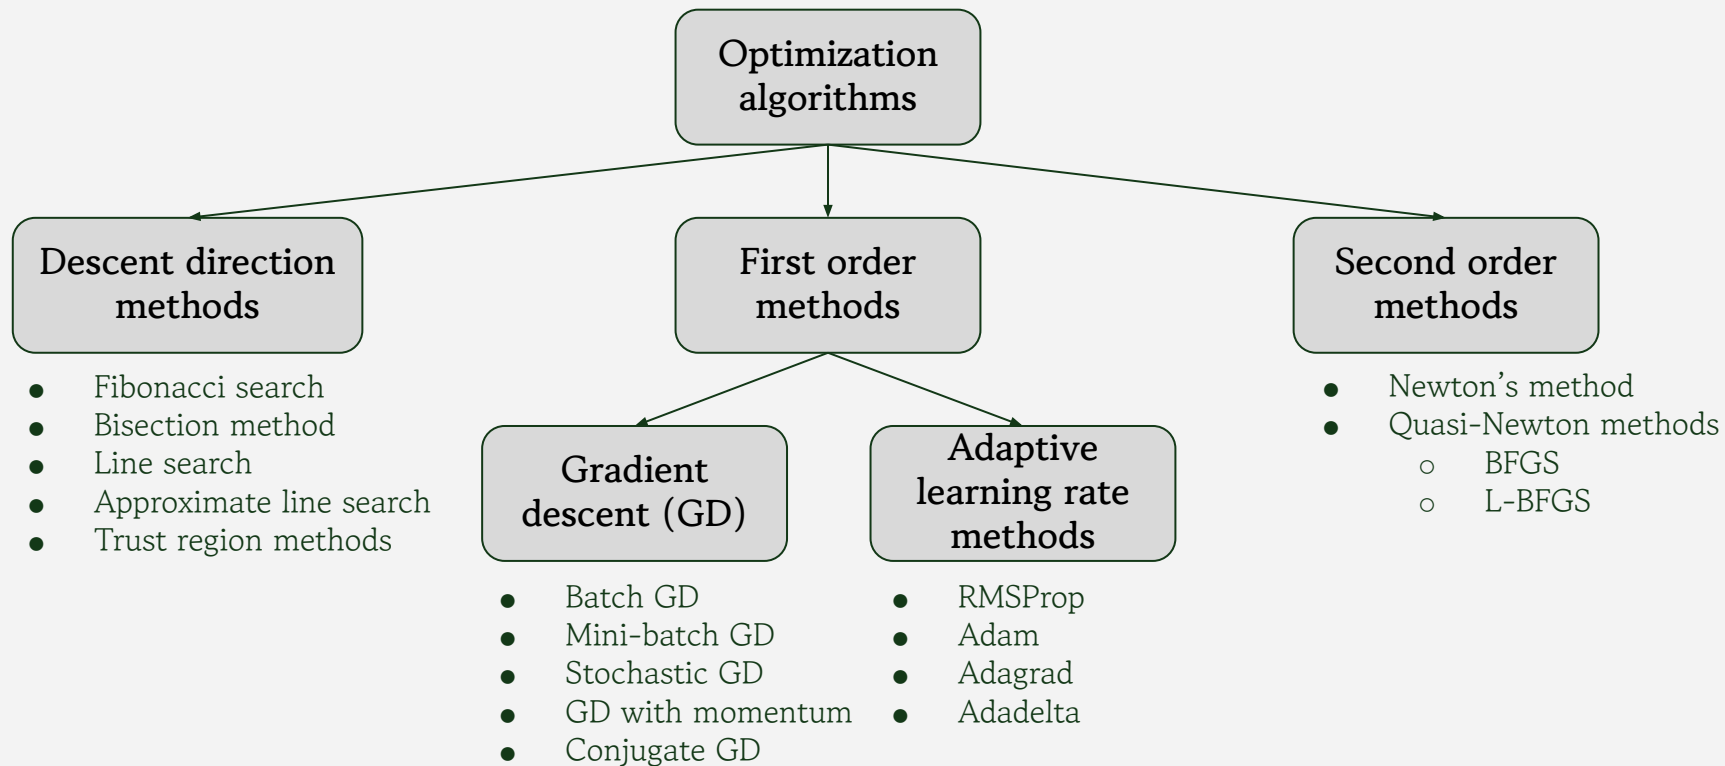
Convergence  
analysis

Codes

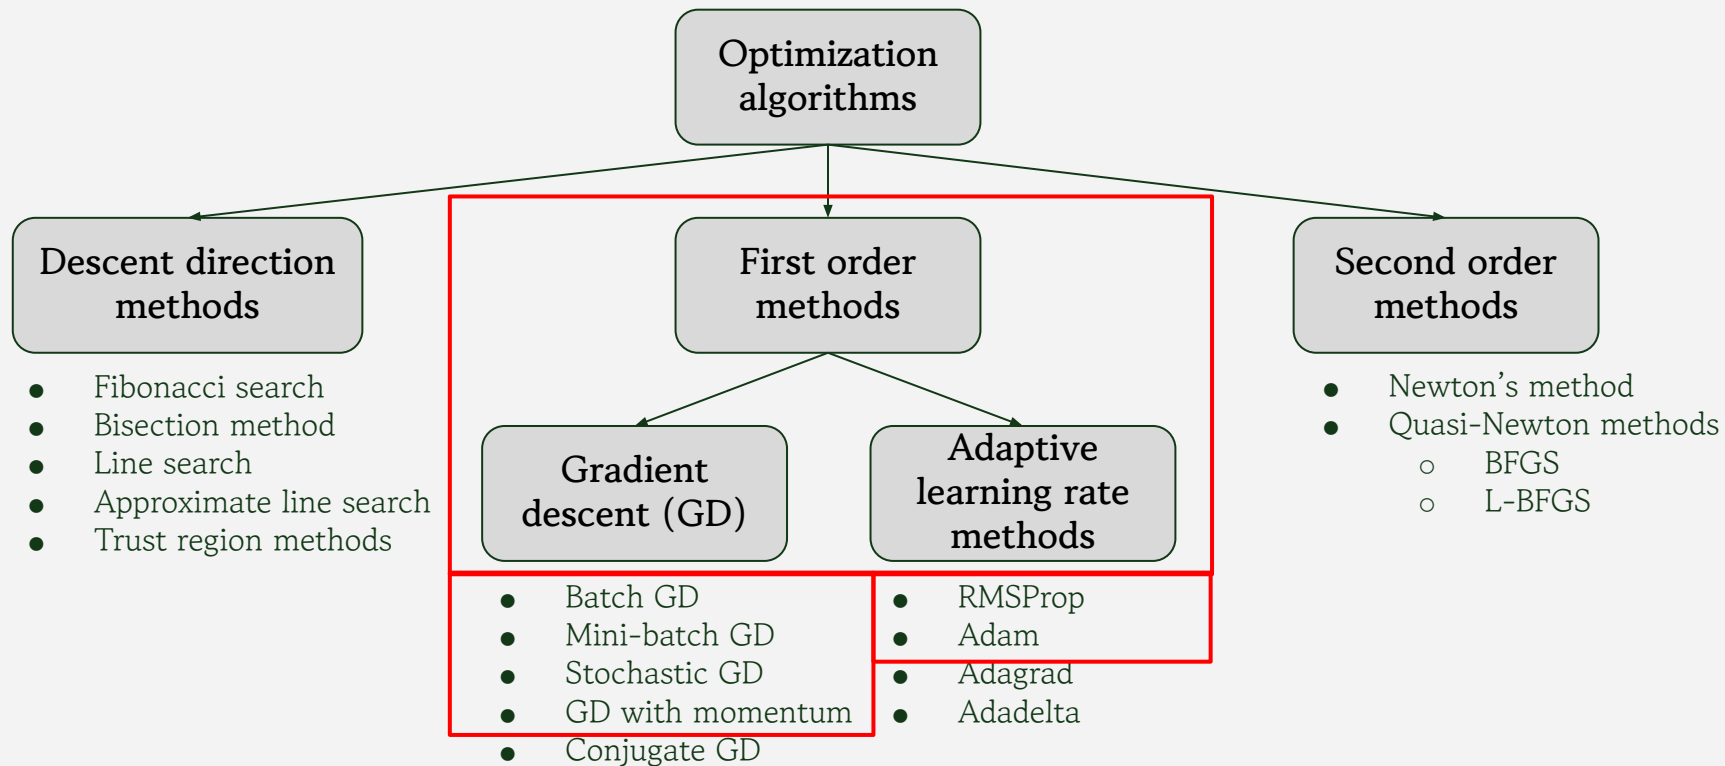
Experiments

Limitations

# Big picture!



# Big picture!



# Line-up

1

2

3

4

5

6

7

Optimization  
algorithms

Related  
work

Adam

Convergence  
analysis

Codes

Experiments


Limitations

# Recap: Gradient descent

$$w = w - \alpha \frac{\partial}{\partial w} [C(w)]$$

$w$ : model parameter

$C(w)$ : objective function



Credits: <https://mlfromscratch.com/optimizers-explained>



## Gradient descent variants

# Batch gradient descent

$$X_{(n \times m)} = [x^{(1)} \ x^{(2)} \ \dots \ x^{(m)}]$$

$$Y_{(1 \times m)} = [y^{(1)} \ y^{(2)} \ \dots \ y^{(m)}]$$

$W$ : weight parameter

$b$ : bias parameter

$J(w, b)$ : loss function

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$$

$$w = w - \alpha \frac{\partial}{\partial w} [J(w, b)]$$

$$b = b - \alpha \frac{\partial}{\partial b} [J(w, b)]$$

- **All training examples** are processed at once
- **High** computational cost (utilizes vectorization w/o for loops)
- **Slow** convergence (1 parameters per epoch)

# Mini-batch gradient descent

$$X_{(n \times m)} = \left[ \underbrace{x^{(1)} \ x^{(2)} \ \dots \ x^{(b)}}_{X_{(n \times b)}^{\{1\}}} \mid \underbrace{x^{(b+1)} \ x^{(b+2)} \ \dots \ x^{(2b)}}_{X_{(n \times b)}^{\{2\}}} \mid \dots \right]$$

$$Y_{(1 \times m)} = \left[ \underbrace{y^{(1)} \ y^{(2)} \ \dots \ y^{(b)}}_{Y_{(1 \times b)}^{\{1\}}} \mid \underbrace{y^{(b+1)} \ y^{(b+2)} \ \dots \ y^{(2b)}}_{Y_{(1 \times b)}^{\{2\}}} \mid \dots \right]$$

$X^{\{t\}}$ :  $t^{\text{th}}$  mini-batch of training examples

$Y^{\{t\}}$ :  $t^{\text{th}}$  mini-batch of targets

$$J^{\{t\}}(w, b) = \frac{1}{b} \sum_{i=1}^b L(\hat{y}^{(i)}, y^{(i)})$$

$$w = w - \alpha \frac{\partial}{\partial w} [J^{\{t\}}(w, b)]$$

$$b = b - \alpha \frac{\partial}{\partial b} [J^{\{t\}}(w, b)]$$

- All training examples are processed in **mini-batches**
- **Medium** computational cost (reduced vectorization w/ for loops)
- **Fastest** learning (parameters get updated  $\text{ceil}(m/b)$  number of steps per epoch)

# Stochastic gradient descent

$$X_{(n \times m)} = \begin{bmatrix} \underbrace{x^{(1)}}_{X_{(n \times 1)}^{\{1\}}} & x^{(2)} & \dots & \underbrace{x^{(m)}}_{X_{(n \times 1)}^{\{m\}}} \end{bmatrix}$$

$$Y_{(1 \times m)} = \begin{bmatrix} \underbrace{y^{(1)}}_{Y_{(1 \times 1)}^{\{1\}}} & y^{(2)} & \dots & \underbrace{y^{(m)}}_{Y_{(1 \times 1)}^{\{m\}}} \end{bmatrix}$$

$b = 1$ : here mini-batch size is one

$X^{\{t\}}$ :  $t^{\text{th}}$  mini-batch of training examples

$Y^{\{t\}}$ :  $t^{\text{th}}$  mini-batch of targets

$$J^{\{t\}}(w, b) = \frac{1}{b} \sum_{i=1}^b L(\hat{y}^{(i)}, y^{(i)})$$


$$w = w - \alpha \frac{\partial}{\partial w} [J^{\{t\}}(w, b)]$$

$$b = b - \alpha \frac{\partial}{\partial b} [J^{\{t\}}(w, b)]$$


- All training examples are processed **one sample** at a time
- **Low** computational cost per update
- **Fast but oscillates**, loses speedup from vectorization (parameters get updated per sample per epoch)

# Training

Low noise in updates



More noise in updates  
(less noise than SGD)



Credits: Deep learning specialization by Prof. Andrew Ng, Coursera

# Stochasticity

Randomness or unpredictability in a system or process

E.g.,

- Mini-batch gradient descent
- SGD



**Stochastic** objective functions

# EWA (Exponentially Weighted Average)

$$v_t = \beta v_{t-1} + (1 - \beta)\theta_t \quad \beta \text{ is usually a decimal close to 1 (e.g. 0.9, 0.99, etc.)}$$

$$v_0 = 0$$

$$v_1 = (1 - \beta)\theta_1$$


$$v_2 = \beta v_1 + (1 - \beta)\theta_2 \iff v_2 = \beta((1 - \beta)\theta_1) + (1 - \beta)\theta_2$$

$$v_3 = \beta v_2 + (1 - \beta)\theta_3 \iff v_3 = \beta(\beta((1 - \beta)\theta_1) + (1 - \beta)\theta_2) + (1 - \beta)\theta_3$$



$$v_3 = \beta^2(1 - \beta)\theta_1 + \beta(1 - \beta)\theta_2 + (1 - \beta)\theta_3$$

# Why 'exponential'?

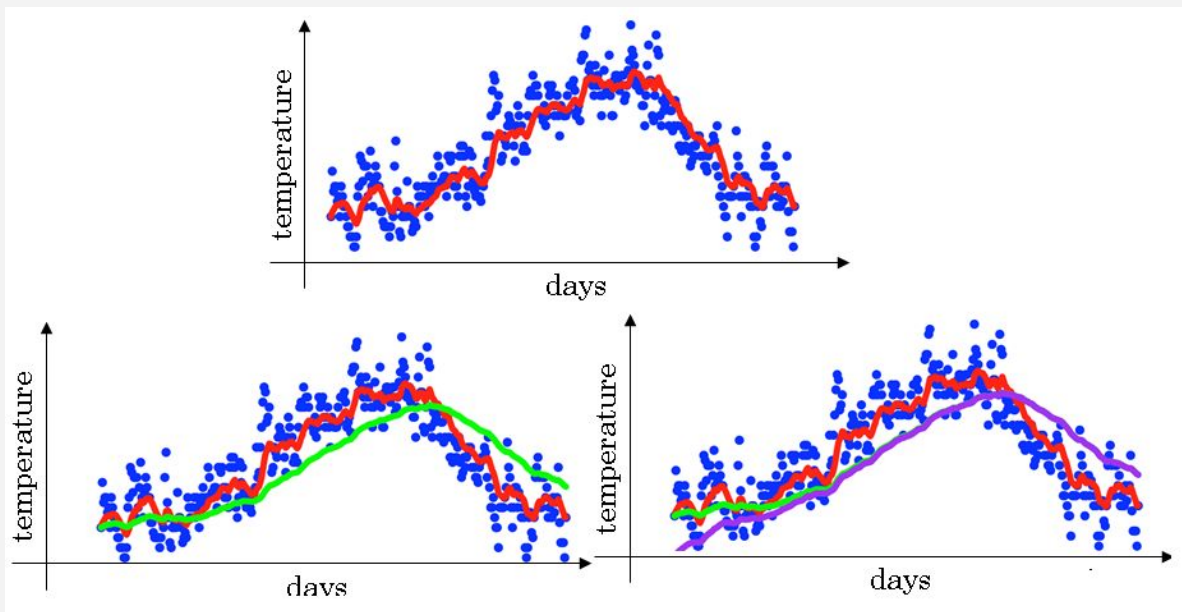


Credits: upscfever.com

$$\beta = 0.9$$
$$(1-\beta) = 0.1$$

- Places greater weights on the most recent data
- The weighting for older data points decrease exponentially
- Computationally efficient
- Saves storage (overrides)

# Bias correction



$$v_t = \beta v_{t-1} + (1 - \beta)\theta_t$$

- Initially biased towards initialization
- Bias corrected term:

$$\hat{v}_t = \frac{v_t}{1 - \beta^t}$$

- **Green plot:** bias not corrected
- **Purple plot:** bias corrected

Credits: Deep learning specialization by Prof. Andrew Ng, Coursera



# Gradient descent with momentum

**First moment** estimate/estimate for **mean of gradients**

$$m_{t,w} = \beta_1 m_{t-1,w} + (1 - \beta_1) g_{t,w}$$

$$m_{t,b} = \beta_1 m_{t-1,b} + (1 - \beta_1) g_{t,b}$$

$$\text{Bias correction: } \hat{m}_{t,w} = \frac{m_{t,w}}{1 - \beta_1^t} \quad \hat{m}_{t,b} = \frac{m_{t,b}}{1 - \beta_1^t}$$


Parameter updates:  $W = W - \alpha \hat{m}_{t,w} \quad b = b - \alpha \hat{m}_{t,b}$

Gradient w.r.t. Stochastic objective at timestep t:

$$g_{t,w} = \frac{\partial}{\partial w} [J^{\{t\}}(w, b)]$$

$$g_{t,b} = \frac{\partial}{\partial b} [J^{\{t\}}(w, b)]$$

- Minimizes the vertical direction gradient component



Credits: andreaperlato.com

# Adagrad (Adaptive gradient)

$$g_t = \nabla_{\theta} J(\theta_t)$$

Accumulates the squares of past gradients:

$$G_t = G_{t-1} + g_t^2$$

Update rule for parameters:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{G_t} + \epsilon} \odot g_t$$

## Adaptive learning rate methods

# RMSPProp

**Second moment** estimate/estimate for **variance of gradients**

$$v_{t,w} = \beta_2 v_{t-1,w} + (1 - \beta_2) g_{t,w}^2$$

$$v_{t,b} = \beta_2 v_{t-1,b} + (1 - \beta_2) g_{t,b}^2$$

Bias correction:

$$\hat{v}_{t,w} = \frac{v_{t,w}}{1 - \beta_2^t} \quad \hat{v}_{t,b} = \frac{v_{t,b}}{1 - \beta_2^t}$$

Parameter updates:

$$W = W - \alpha \frac{g_{t,w}}{\sqrt{\hat{v}_{t,w} + \epsilon}}$$

$$b = b - \alpha \frac{g_{t,b}}{\sqrt{\hat{v}_{t,b} + \epsilon}}$$

Root  
Mean  
Square  
Propagation

Neural Networks for Machine Learning

Lecture 6e

rmsprop: Divide the gradient by a running average of its recent magnitude

Geoffrey Hinton  
with  
Nitish Srivastava  
Kevin Swersky



- Minimizes variance in gradients

RMSPProp

# Line-up

1

2

3

4

5

6

7

Optimization  
algorithms

Related  
work

Adam

Convergence  
analysis

Codes

Experiments

Limitations

# Adam (Adaptive moment estimation)

1st moment estimate (momentum term):

$$m_{t,w} = \beta_1 m_{t-1,w} + (1 - \beta_1) g_{t,w}$$

$$\hat{m}_{t,w} = \frac{m_{t,w}}{1 - \beta_1^t}$$

$$m_{t,b} = \beta_1 m_{t-1,b} + (1 - \beta_1) g_{t,b}$$

$$\hat{m}_{t,b} = \frac{m_{t,b}}{1 - \beta_1^t}$$

2nd moment estimate (RMSProp term):

$$v_{t,w} = \beta_2 v_{t-1,w} + (1 - \beta_2) g_{t,w}^2$$

$$\hat{v}_{t,w} = \frac{v_{t,w}}{1 - \beta_2^t}$$

$$v_{t,b} = \beta_2 v_{t-1,b} + (1 - \beta_2) g_{t,b}^2$$

$$\hat{v}_{t,b} = \frac{v_{t,b}}{1 - \beta_2^t}$$

Parameter updates:

$$W = W - \alpha \frac{\hat{m}_{t,w}}{\sqrt{\hat{v}_{t,w} + \varepsilon}}$$

$$b = b - \alpha \frac{\hat{m}_{t,b}}{\sqrt{\hat{v}_{t,b} + \varepsilon}}$$

# Algorithm

**$\beta_1$** : Exponential decay rate parameter for past gradient estimates

**$\beta_2$** : Exponential decay rate parameter for past squared gradient estimates

---

**Algorithm 1:** *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation.  $g_t^2$  indicates the elementwise square  $g_t \odot g_t$ . Good default settings for the tested machine learning problems are  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$  and  $\epsilon = 10^{-8}$ . All operations on vectors are element-wise. With  $\beta_1^t$  and  $\beta_2^t$  we denote  $\beta_1$  and  $\beta_2$  to the power  $t$ .

---

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1)$ : Exponential decay rates for the moment estimates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)

$v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)

**end while**

**return**  $\theta_t$  (Resulting parameters)

---

# Adam's update rule

$$\Delta_t = \alpha \cdot \hat{m}_t / \sqrt{\hat{v}_t}$$

$$\text{case } (1 - \beta_1) > \sqrt{1 - \beta_2} \quad |\Delta_t| \leq \alpha \cdot (1 - \beta_1) / \sqrt{1 - \beta_2}$$

If gradient has been zero at all time steps except current step,

$$\left. \begin{aligned} m_t &\approx (1 - \beta_1) g_t \Rightarrow \hat{m}_t \approx \frac{(1 - \beta_1) g_t}{(1 - \beta_1^t)} \\ v_t &\approx (1 - \beta_2) g_t^2 \Rightarrow \hat{v}_t \approx \frac{(1 - \beta_2) g_t^2}{(1 - \beta_2^t)} \end{aligned} \right\} \text{when } t \uparrow \beta_1^t, \beta_2^t \simeq 0$$

$$\begin{aligned} \hat{m}_t &\simeq m_t \\ \hat{v}_t &\simeq v_t \end{aligned} \Rightarrow \left| \frac{\hat{m}_t}{\sqrt{\hat{v}_t}} \right| = \left| \frac{(1 - \beta_1) g_t}{\sqrt{(1 - \beta_2) g_t^2}} \right| = \frac{1 - \beta_1}{\sqrt{1 - \beta_2}}$$

$$\begin{aligned} (1 - \beta_1) &> \sqrt{1 - \beta_2} \\ |\Delta_t| &\leq \alpha \cdot \frac{(1 - \beta_1)}{\sqrt{1 - \beta_2}} \end{aligned}$$

$\prec \alpha \cdot 1$

Tighter upper bound (for sparsity)

# Adam's update rule

$$\Delta_t = \alpha \cdot \hat{m}_t / \sqrt{\hat{v}_t}$$

Common case:  $|\Delta_t| \leq \alpha$

SNR (signal to noise ratio)

Jensen's inequality

For any stochastic gradient,  $g_t$ :

1st moment estimate  $\hat{m}_t \rightarrow \frac{E[g_t]}{\sqrt{E[g_t^2]}} \leq 1$

2nd moment estimate  $\hat{v}_t$

$\therefore$  General case  $|\Delta_t| \leq \alpha \cdot 1$

- Adam is bounded
- Does not grow exponentially



# Line-up

1

2

3

4

5

6

7

Optimization  
algorithms

Related  
work

Adam

Convergence  
analysis

Codes

Experiments

Limitations

# Line-up

1

2

3

4

5

6

7

Optimization  
algorithms

Related  
work

Adam

Convergence  
analysis

Codes

Experiments

Limitations

# PyTorch

## Adam

```
CLASS torch.optim.Adam(params, lr=0.001, betas=(0.9, 0.999), eps=1e-08, weight_decay=0,  
    amsgrad=False, *, foreach=None, maximize=False, capturable=False, differentiable=False,  
    fused=None) [SOURCE]
```

### Parameters

- **params** (*iterable*) – iterable of parameters or named\_parameters to optimize or iterable of dicts defining parameter groups. When using named\_parameters, all parameters in all groups should be named
- **lr** (*float*, *Tensor*, *optional*) – learning rate (default: 1e-3). A tensor LR is not yet supported for all our implementations. Please use a float LR if you are not also specifying fused=True or capturable=True.
- **betas** (*Tuple[*float*, *float*]*, *optional*) – coefficients used for computing running averages of gradient and its square (default: (0.9, 0.999))
- **eps** (*float*, *optional*) – term added to the denominator to improve numerical stability (default: 1e-8)
- **weight\_decay** (*float*, *optional*) – weight decay (L2 penalty) (default: 0)

# Line-up

1

2

3

4

5

6

7

Optimization  
algorithms

Related  
work

Adam

Convergence  
analysis

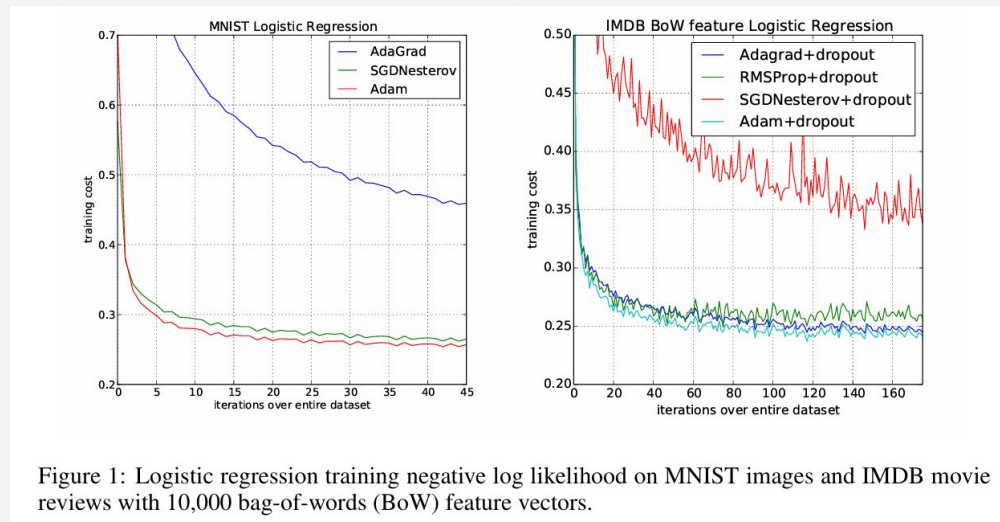
Codes

Experiments

Limitations

# Experiments

## Logistic regression (convex)



Adam's learning rate decay:

$$\alpha_t = \frac{\alpha}{\sqrt{t}}$$

- The 10,000 dimension BoW feature vector for each review is highly sparse
- Adam can take advantage of sparse features and obtain faster convergence

# Experiments

## Multi-layer neural network (non-convex)

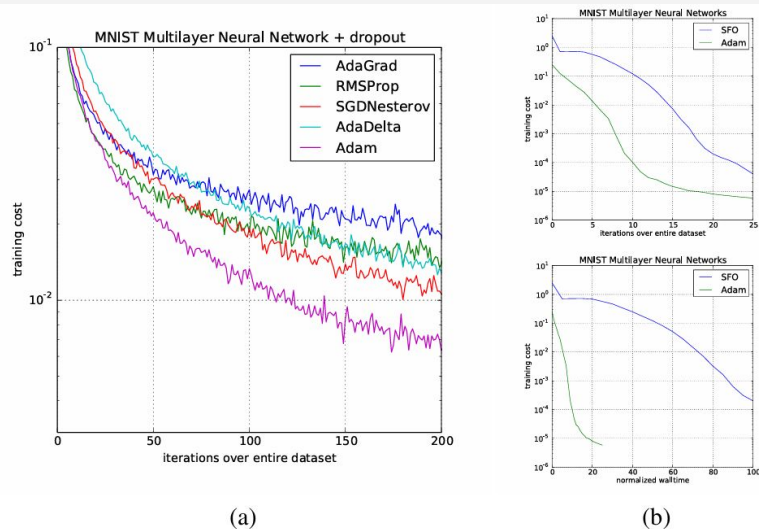


Figure 2: Training of multilayer neural networks on MNIST images. (a) Neural networks using dropout stochastic regularization. (b) Neural networks with deterministic cost function. We compare with the sum-of-functions (SFO) optimizer (Sohl-Dickstein et al., 2014)

- A neural network model with:
  - 2 fully connected
  - Hidden layers with 1000 hidden units each
  - ReLU activation
- Mini-batch size of 128
- L2 weight decay on the parameters to prevent over-fitting
- Normalized walltime: actual elapsed time (walltime) adjusted or normalized relative to a reference standard

# Experiments

## CNN (non-convex)

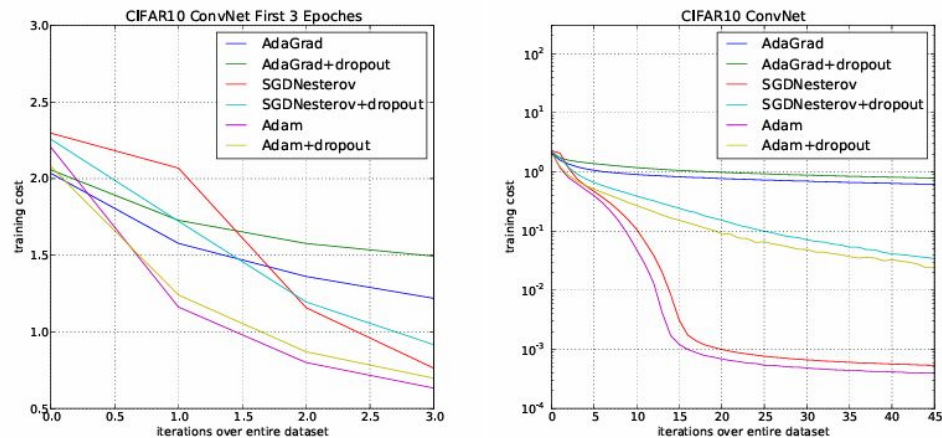


Figure 3: Convolutional neural networks training cost. (left) Training cost for the first three epochs. (right) Training cost over 45 epochs. CIFAR-10 with c64-c64-c128-1000 architecture.

- CNN architecture
  - three alternating stages of 5x5 convolution filters
  - 3x3 max pooling with stride of 2
  - fully connected layer of 1000 rectified linear hidden units (ReLU's)

# Experiments (Ablation)

## Bias correction term

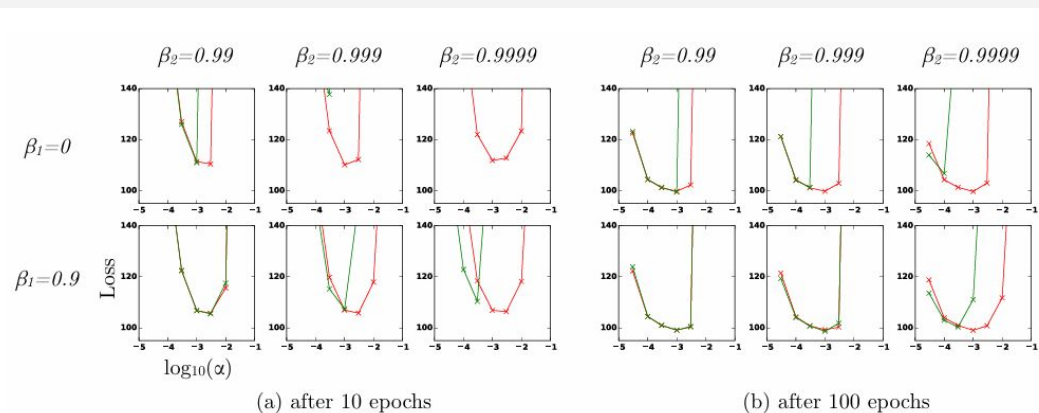


Figure 4: Effect of bias-correction terms (red line) versus no bias correction terms (green line) after 10 epochs (left) and 100 epochs (right) on the loss (y-axes) when learning a Variational Auto-Encoder (VAE) (Kingma & Welling, 2013), for different settings of stepsize  $\alpha$  (x-axes) and hyper-parameters  $\beta_1$  and  $\beta_2$ .

- Values of  $\beta_2$  close to 1, required for robustness to sparse gradients, results in larger initialization bias
- values  $\beta_2$  close to 1 indeed lead to instabilities in training when no bias correction term was present, especially at first few epochs of the training



# Extensions

## AdaMax

**Algorithm 2:** *AdaMax*, a variant of Adam based on the infinity norm. See section 7.1 for details. Good default settings for the tested machine learning problems are  $\alpha = 0.002$ ,  $\beta_1 = 0.9$  and  $\beta_2 = 0.999$ . With  $\beta_1^t$  we denote  $\beta_1$  to the power  $t$ . Here,  $(\alpha/(1 - \beta_1^t))$  is the learning rate with the bias-correction term for the first moment. All operations on vectors are element-wise.

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1)$ : Exponential decay rates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)

$u_0 \leftarrow 0$  (Initialize the exponentially weighted infinity norm)

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$u_t \leftarrow \max(\beta_2 \cdot u_{t-1}, |g_t|)$  (Update the exponentially weighted infinity norm)

$\theta_t \leftarrow \theta_{t-1} - (\alpha/(1 - \beta_1^t)) \cdot m_t/u_t$  (Update parameters)

**end while**

**return**  $\theta_t$  (Resulting parameters)

- We don't need to correct for initialization bias in this case
- Also note that the magnitude of parameter updates has a simpler bound with AdaMax than Adam

# Line-up

1

Optimization  
algorithms

2

Related  
work

3

Adam

4

Convergence  
analysis

5

Codes

6

Experiments

7

Limitations

# Limitations

- **Non-Convergence Issues:** fail to converge to global minima in certain scenarios, particularly when the learning rates do not diminish over time

*[Dereich et al. 2024. Non-convergence of Adam and other adaptive stochastic gradient descent optimization methods for non-vanishing learning rates]*

- **Limit Cycles and Oscillations:** Adam can exhibit limit cycles, leading to oscillations around suboptimal points rather than converging smoothly to an optimal solution

*[Bock et al. 2019. Non-convergence and Limit Cycles in the Adam Optimizer]*

# Thank you!



Credits: meme-arsenal